


# Apache Beam




# Agenda

1. Overview of Apache Beam
2. Crash Course
3. Demo

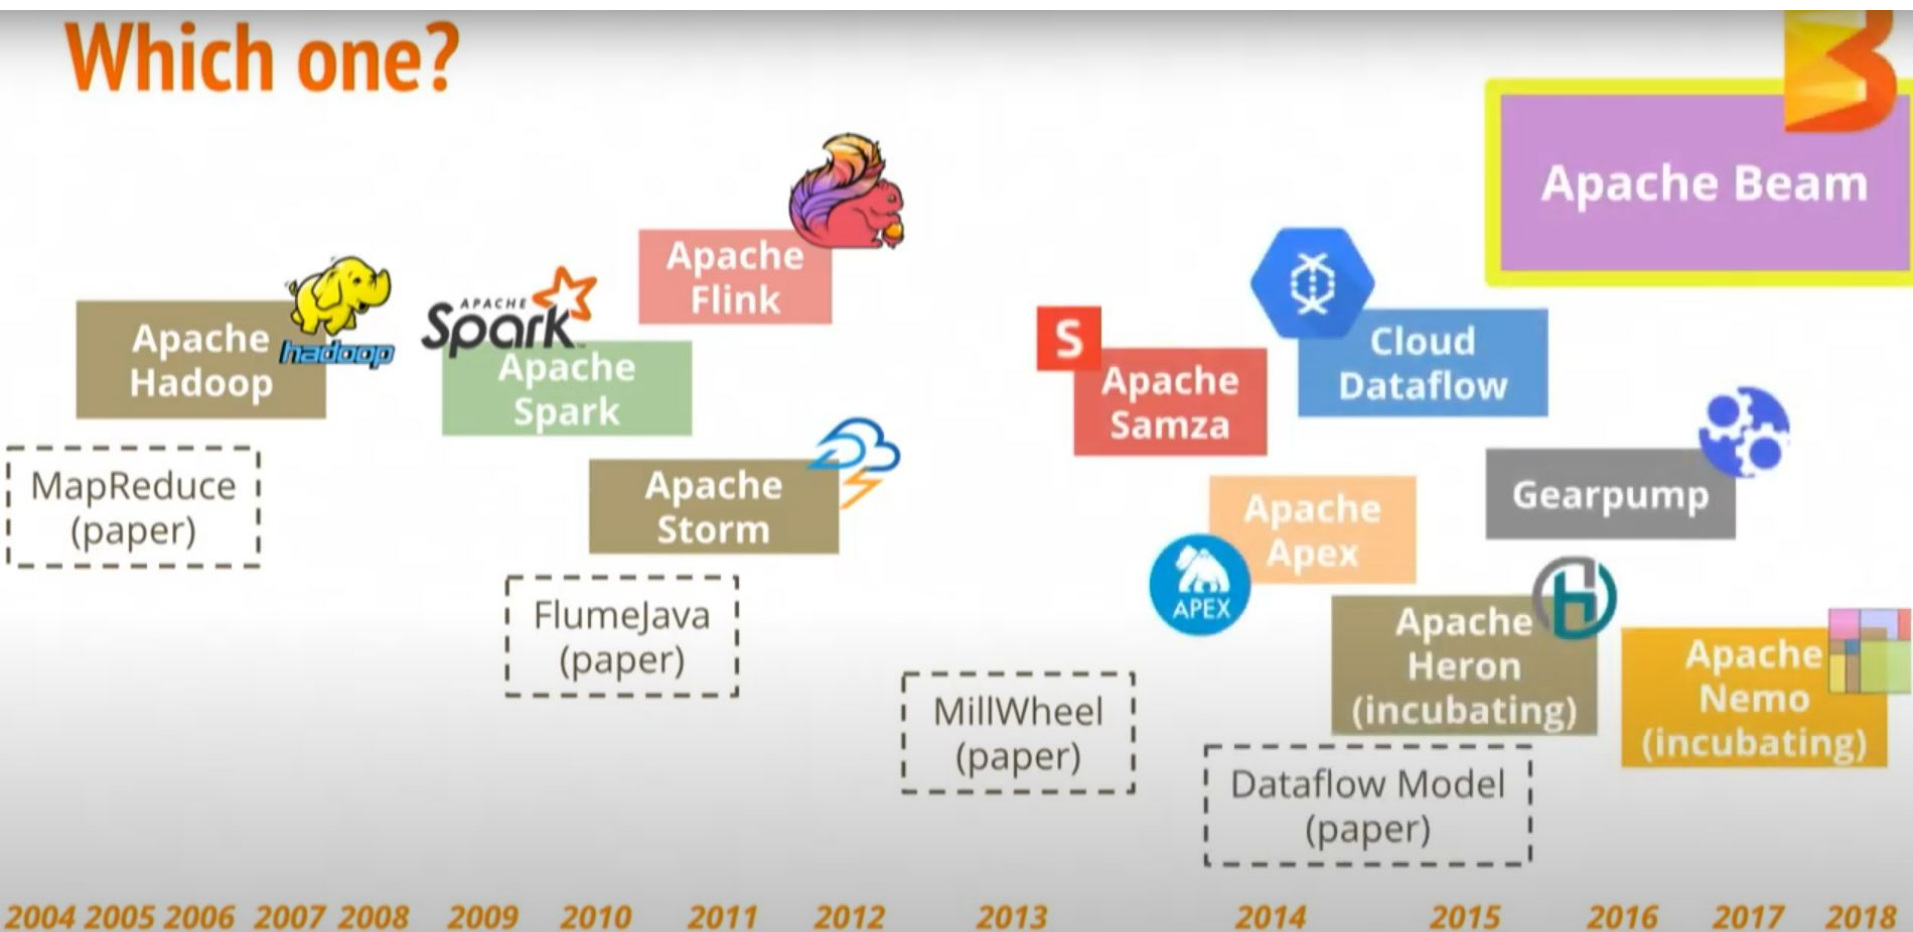
# Scenarios



# Scenarios



# Which one?



# The Beam Vision

Java

```
input.apply(  
    Sum.integersPerKey()  
)
```

Python

```
input | Sum.PerKey()
```

Go

```
stats.Sum(s, input)
```


SQL

```
SELECT key, SUM(value) FROM  
input GROUP BY key
```

Sum Per Key



Cloud Dataflow



Apache Flink



Apache Spark



Apache Apex



Gearpump



IBM Streams




Apache Samza



Apache Nemo  
(incubating)


# Crash Beam - Per element ParDo(Map, etc)



Every Item processed independently.

Stateless implementation.

# Crash Beam - Per key Combine(Reduce, etc)



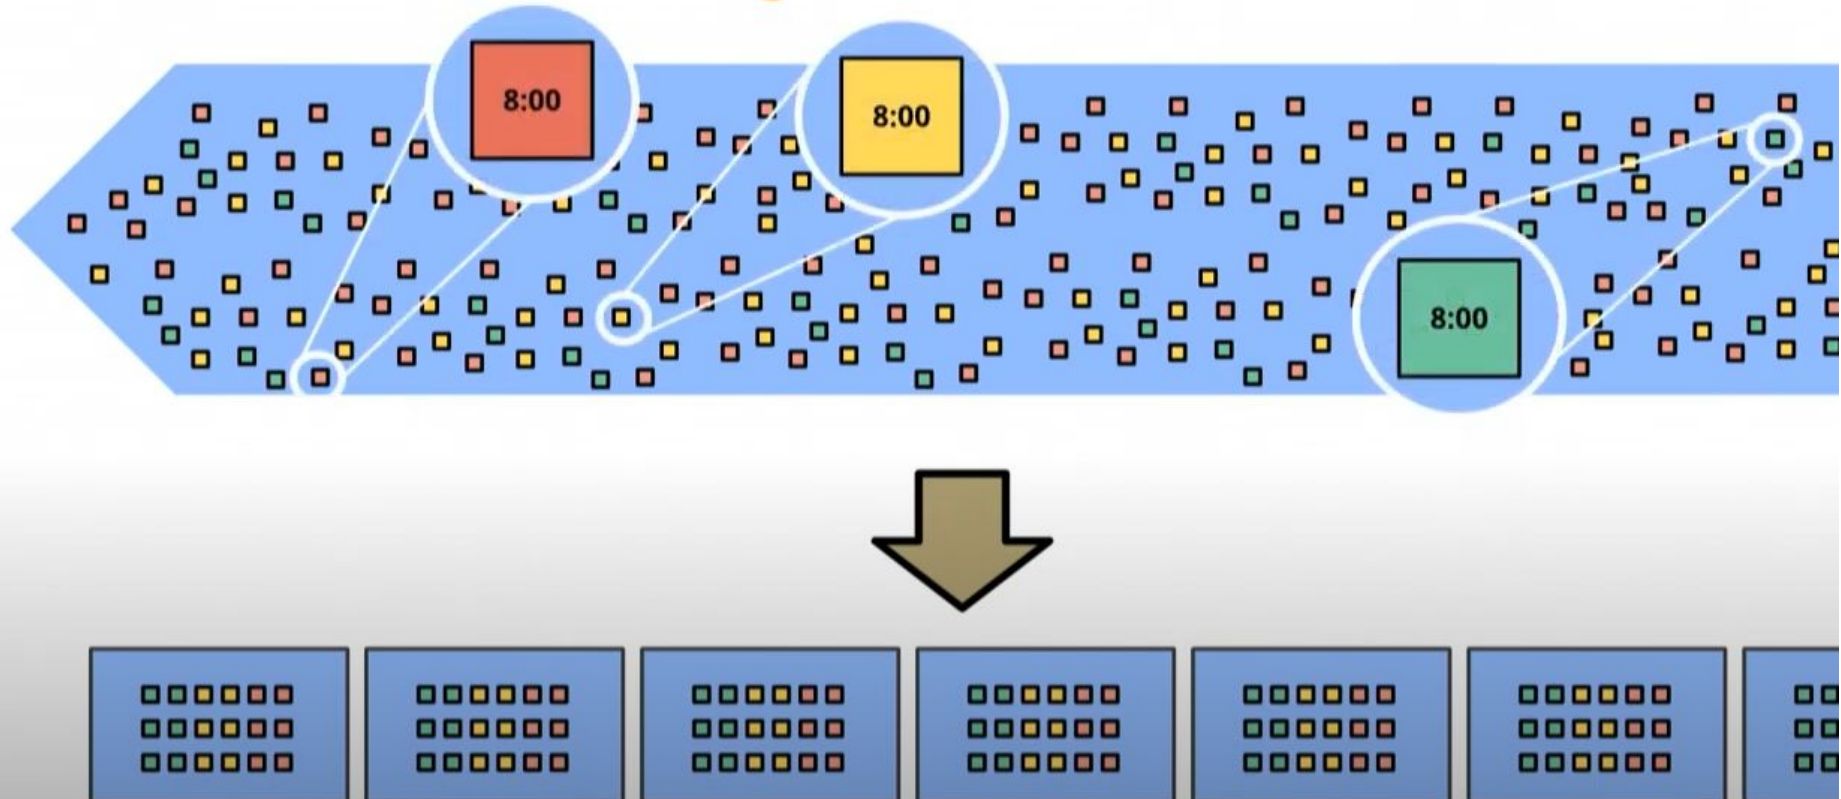
Items grouped by some key and combined

Stateful streaming implementation

But your code doesn't work with state, just associative & commutative function



# Crash Beam - Event Time Windowing



# Crash Summary

1. ParDo: per element processing
2. Combine / GroupByKey: aggregation
3. Event time windowing

# Terms


- **PCollection** — represents a data set which can be a fixed batch or a stream data.
- **PTransform** — a data processing operation that takes one or more PCollections and outputs zero or more PCollections.
- **Pipeline** — represent a directed acyclic graph of PCollection and Transform, and hence encapsulates the entire data processing job.
- **I/O Transforms** — PTransforms that read or write data.

# Demo

Beam currently supports:

Direct runner, Apache Flink runner, Apache Spark runner, Google Cloud Data Flow runner, AWS Kinesis, Apache Nemo runner, Apache Samza runner, Hazelcast Jet runner and Twister2 runner.

A simple linear pipeline with 3 sequential transforms



# References

Apache Beam is a typical Kappa Architecture:

<https://www.peerislands.io/data-processing-lambda-vs-kappa-architectures-and-apache-beam/>