


Kubernetes 101 Tutorial - 0


03/19/2023

A Sample Enterprise System




Ensure maximum availability and scalability

Chat App




Frontoffice

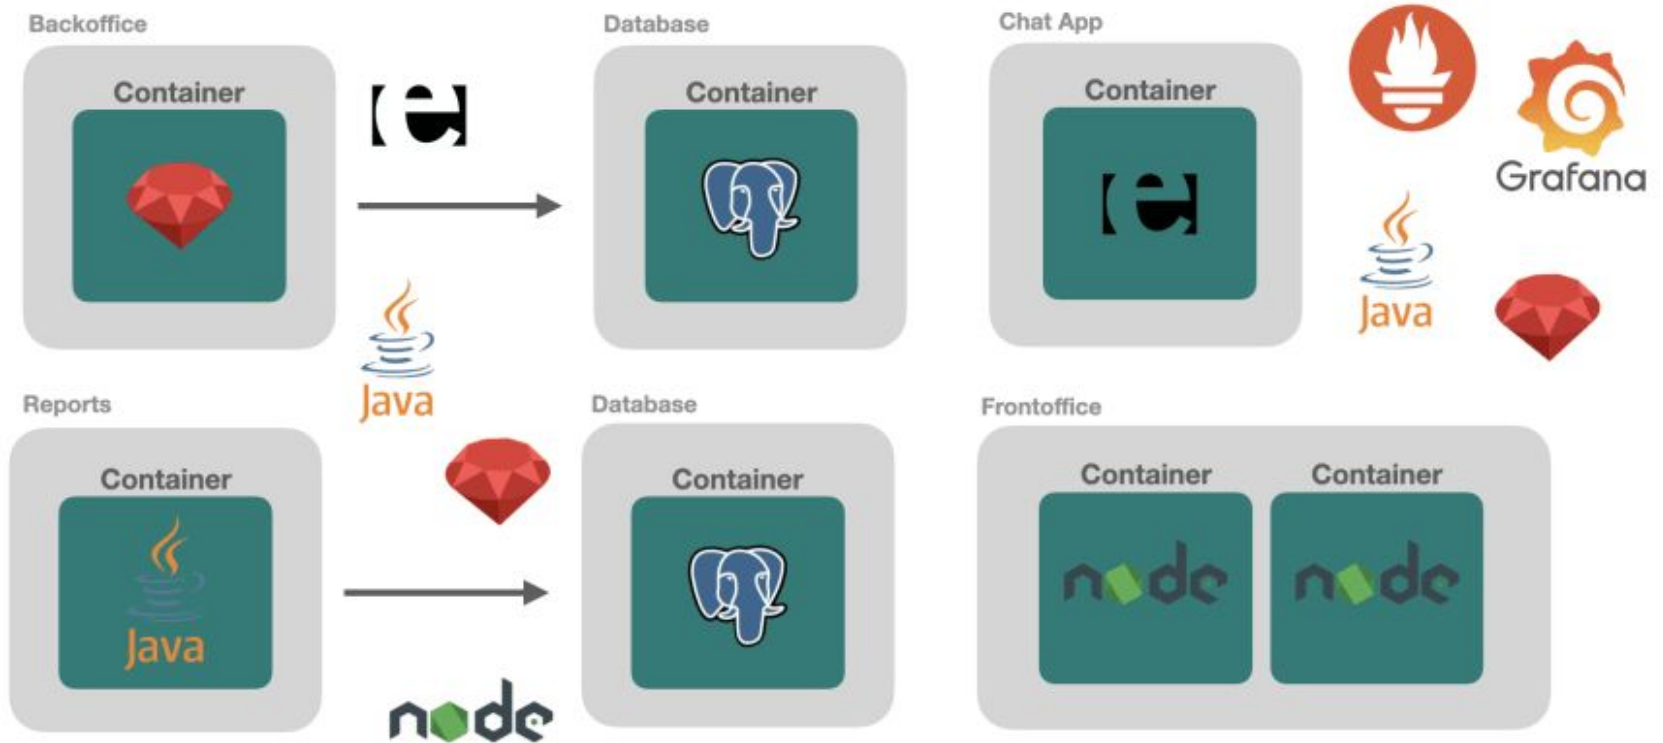


A furthermore functional requirement

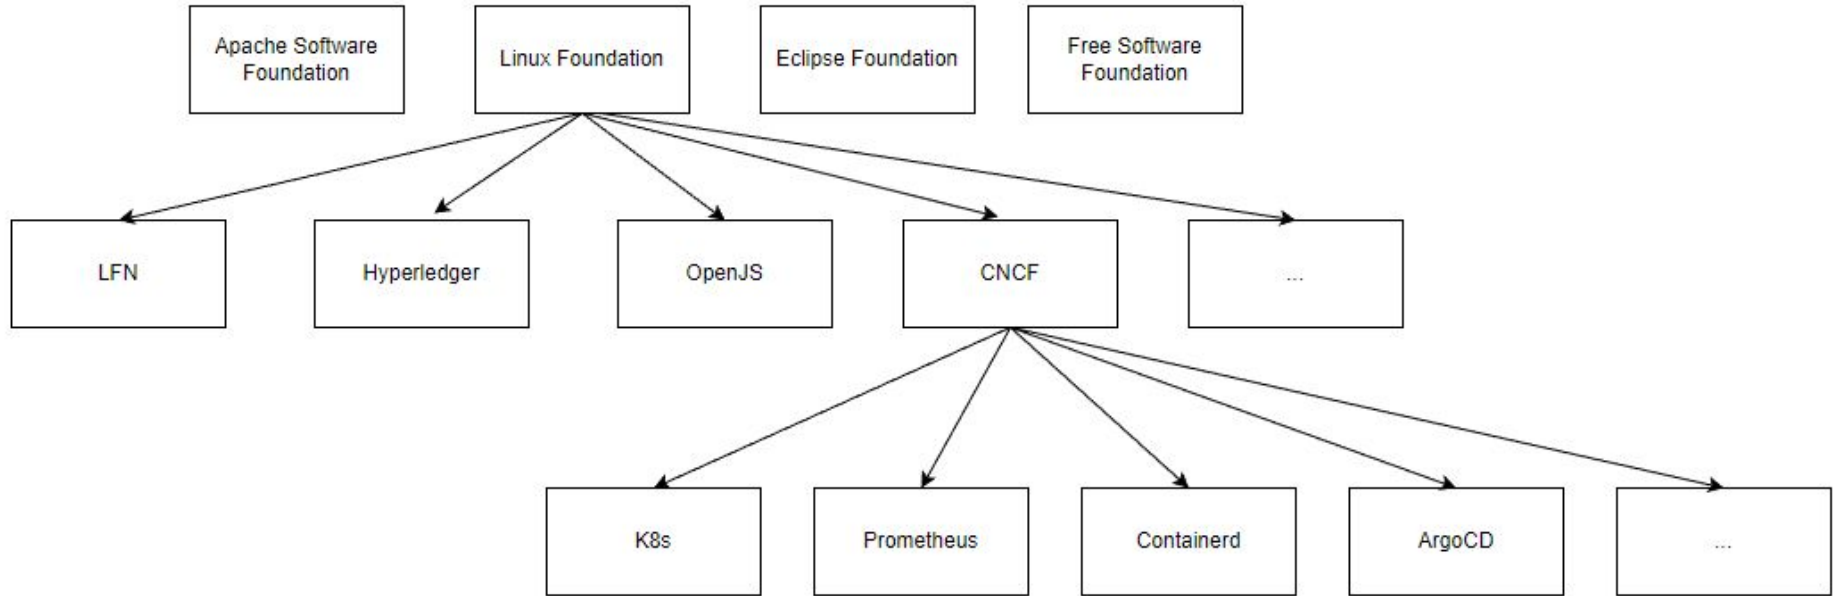
Chat app cannot be down for much time, and in case it goes down, we should **make sure it is started again**, having the capability of **self-healing**



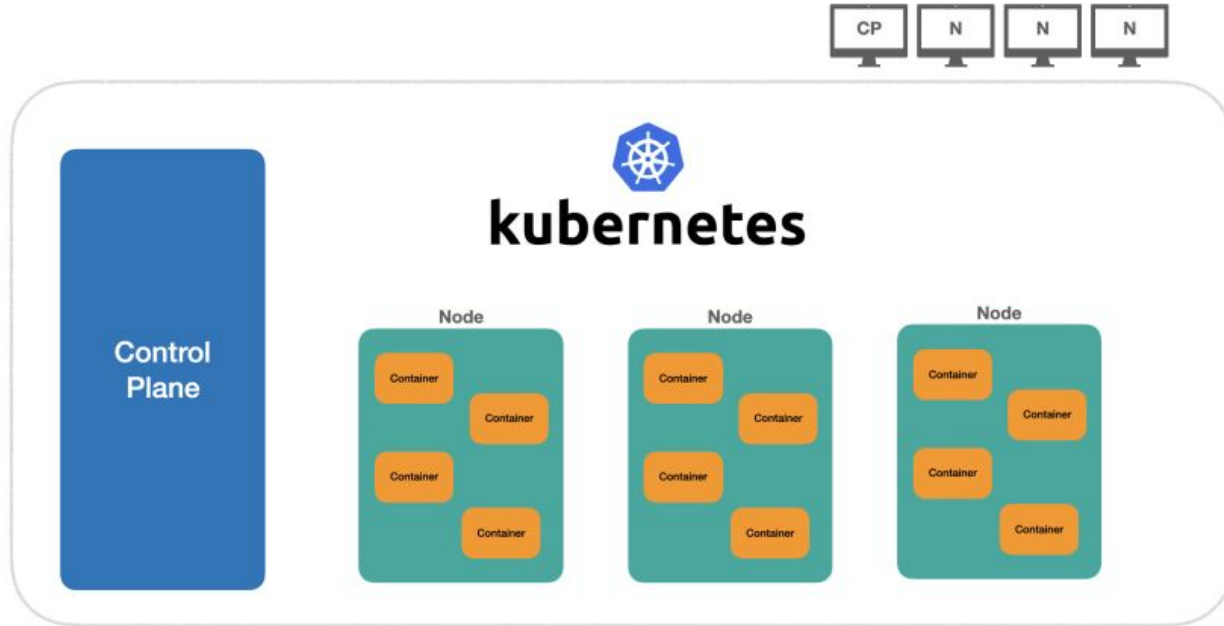
Container Management



A bit history - Linux Foundation Architecture



K8s Architecture



- 1 machine called **Control Plane**, in which the cluster is created and is responsible to accept new machines (or nodes) on the cluster
- 3 other machines called **Nodes**, which will contain all the managed containers by the cluster.

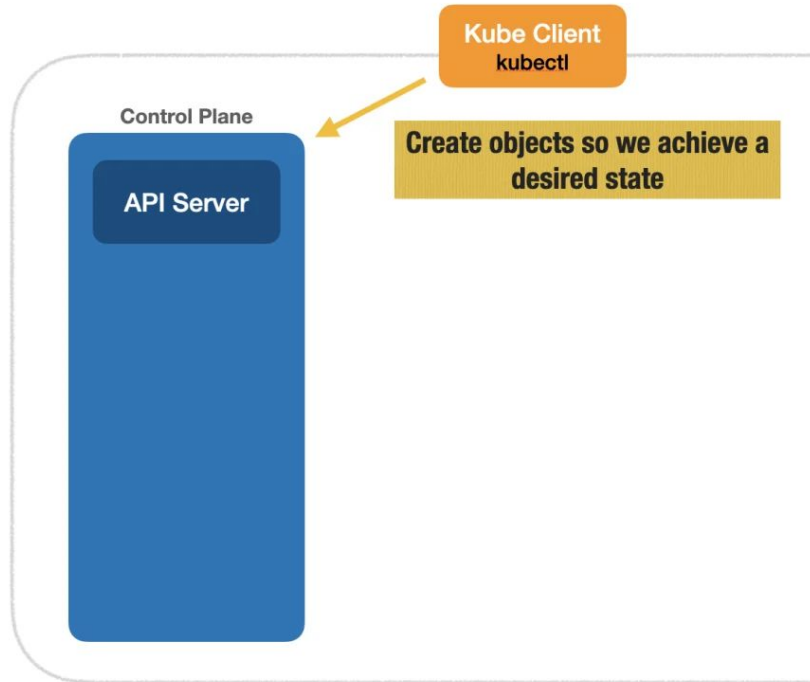
A rule of thumb

All the running containers establish what we call the **cluster state**.

In Kubernetes, we declare the desired state of the cluster by making HTTP requests to the Kubernetes API, and Kubernetes will "work hard" to achieve the **desired state**.

However, making plain HTTP requests in order to declare the state can be somewhat cumbersome, error prone and a tedious job. How about having some CLI in the command-line which would do the hard work of authenticating and making HTTP requests? **Kubect!**

Creating Objects in K8s Cluster



Demo: create an nginx object

What is pod

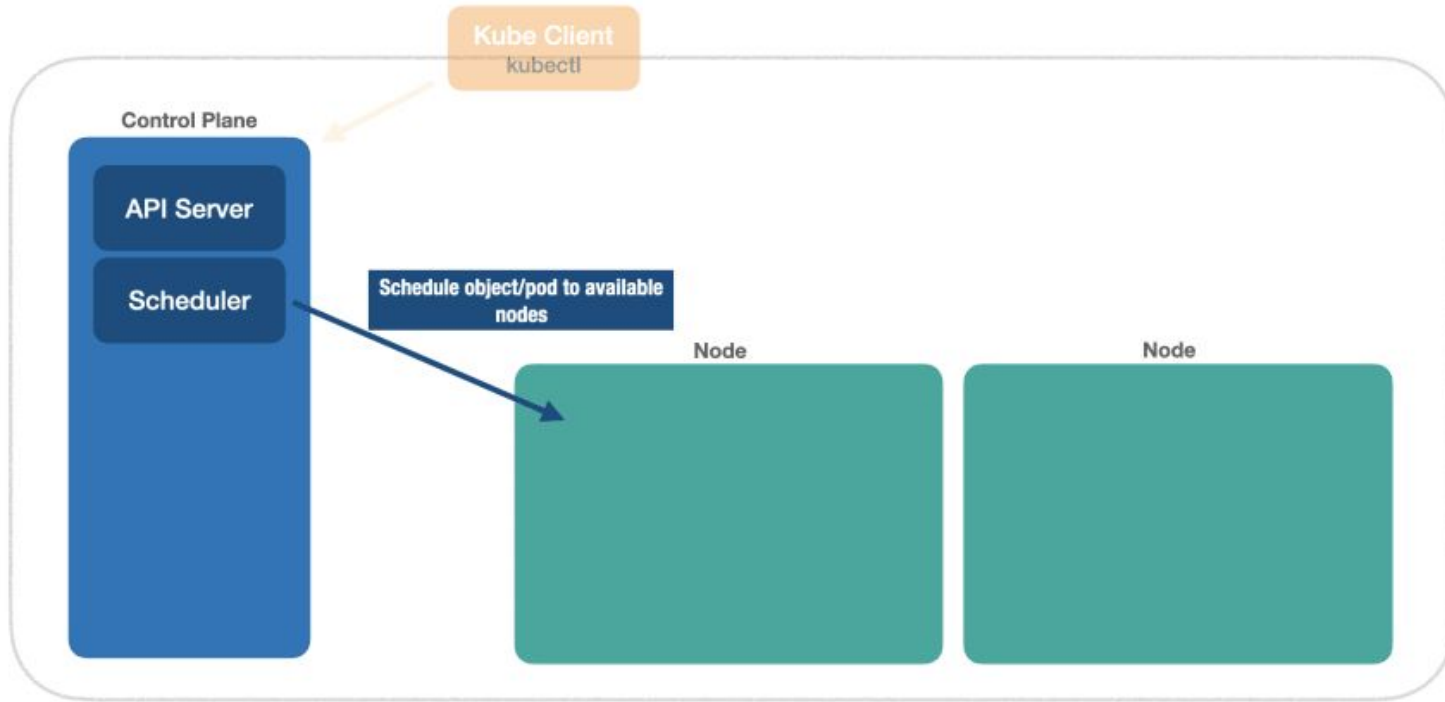
```
1 $ kubectl run nginx --image=nginx  
2 pod/nginx created
```

Pod

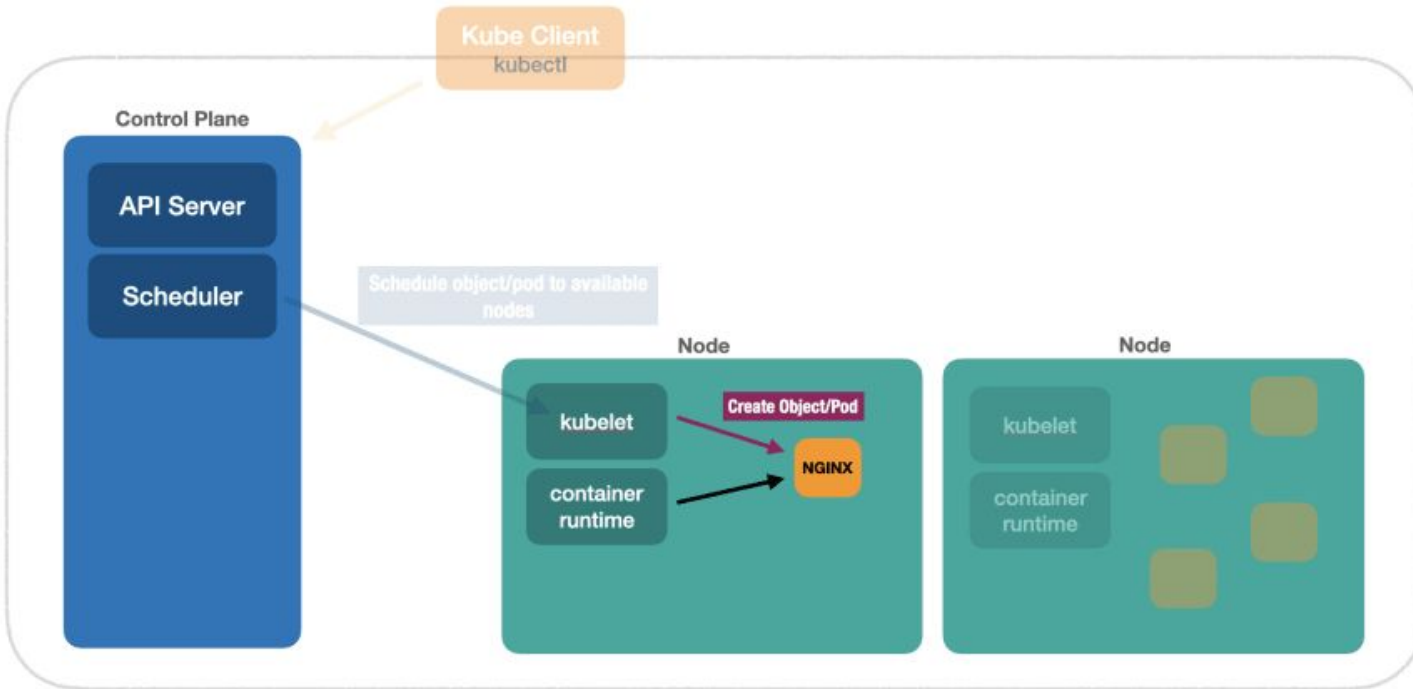
Container

Container

Architecture Flow



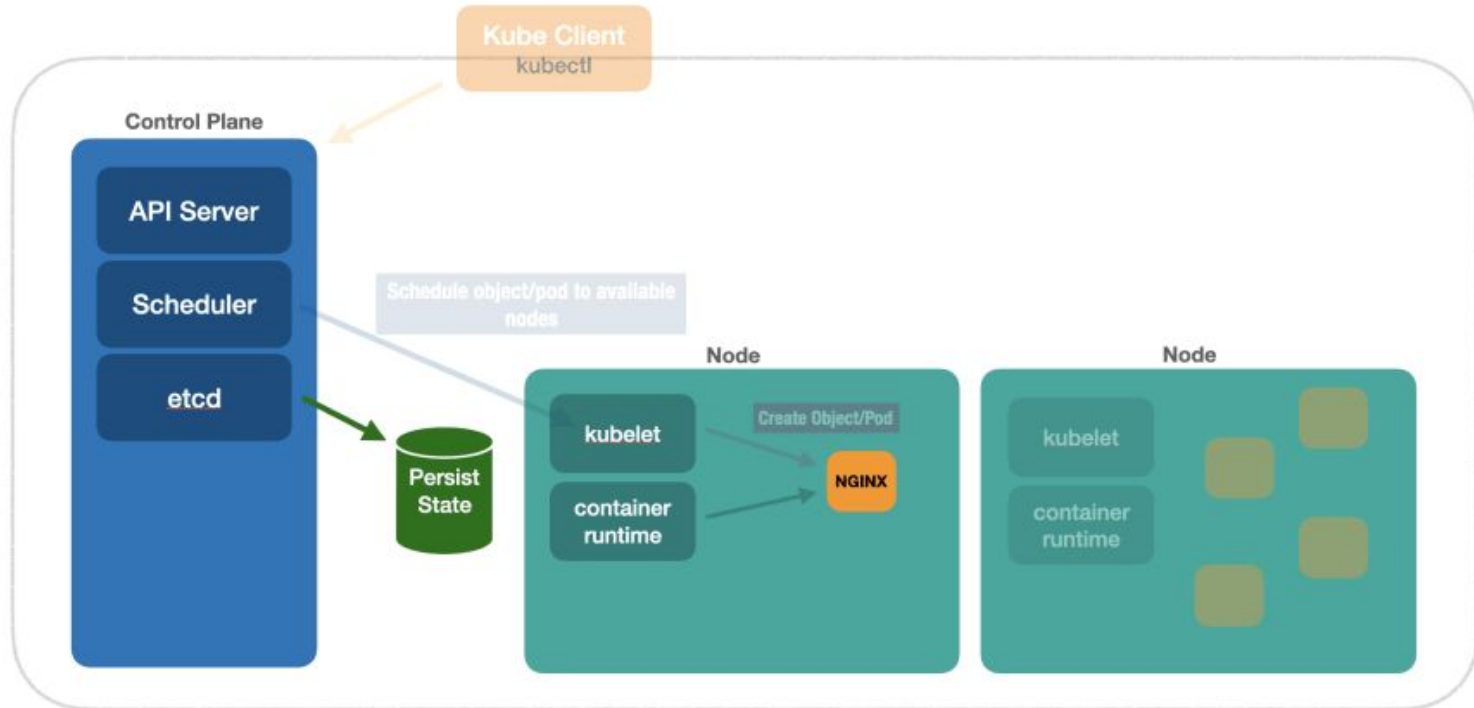
Node kubelet



Each node contains a component called **Kubelet**, which admits objects coming from the Scheduler.

And, using the container runtime installed in the node (could be Docker, containerd, CRI-O, etc), creates the object in the node

etcd - A distributed k-v store



A bit K8s networking

Create 2 deployments:

```
kubectl run server --image=nginx
```

```
kubectl run client --image=nginx
```

In containerized applications, by default, containers are isolated and do not share the host network.

Neither do Pods.

Execute commands in a running pod (server):

```
kubectl exec server -- curl localhost
```

Execute commands in another running pod (client):

```
kubectl describe pod server | grep IP
```

```
kubectl exec client -- curl 172.17.0.6
```

Service

in case we perform a deploy, i.e change the old *server* Pod to a *newer Pod*, **there's no guarantee that the new Pod will get the same previous IP.**

We need some mechanism of **pod discovery**, where we can declare a *special object* in Kubernetes that will **give a name** to a given pod. Therefore, within the cluster, we could **reach Pods by their names** instead of internal IP's.

Controller Manager

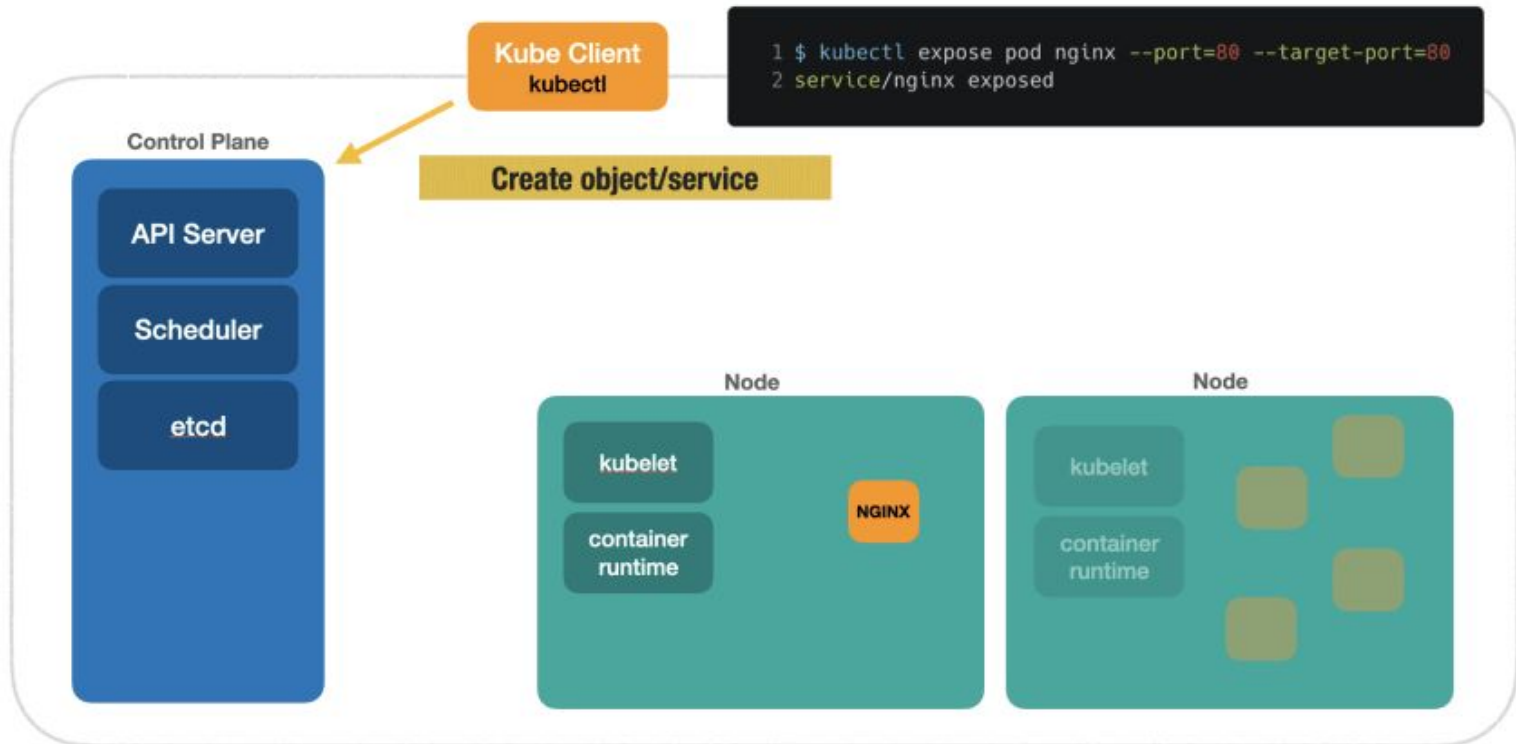
responsible to receive a request for special objects like **Services** and expose them via service discovery.

Expose server name (not IP):

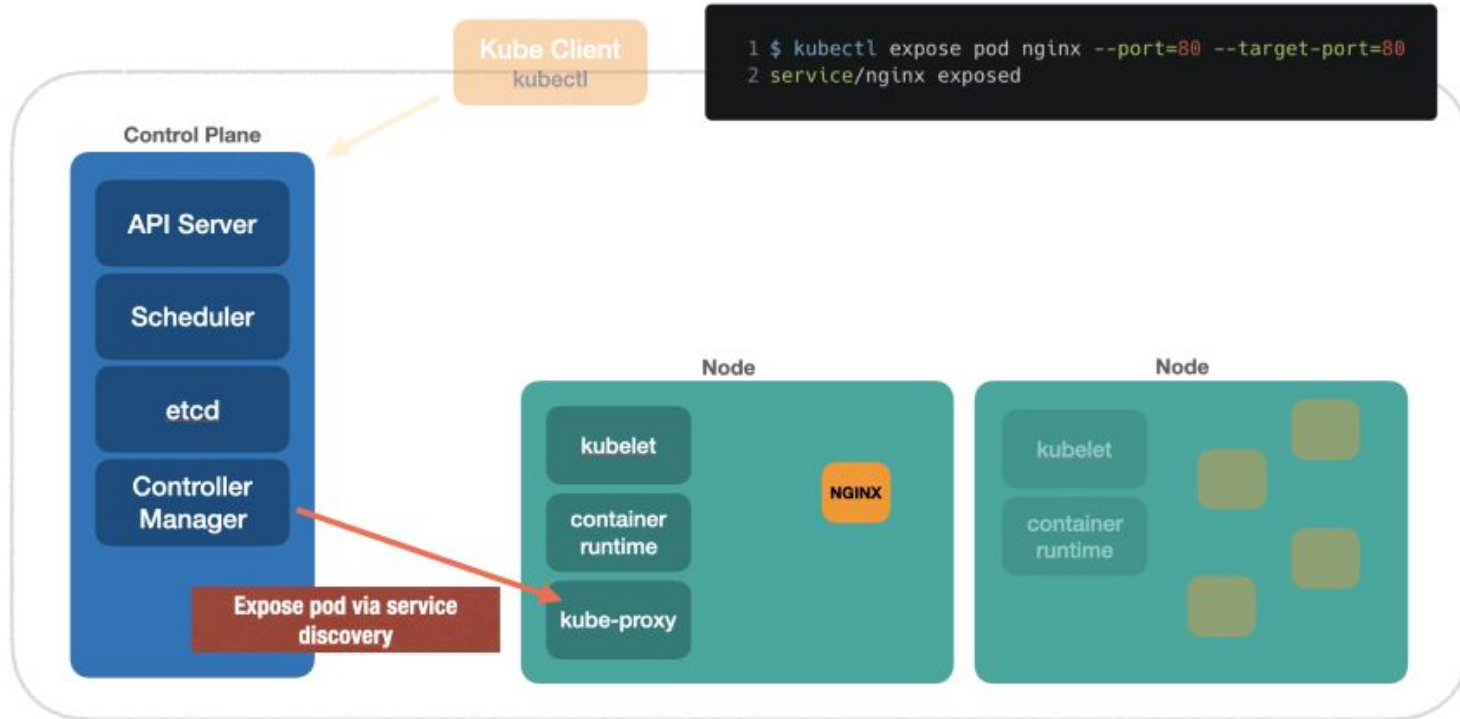
```
kubectll expose pod server --port=80 --target-port=80
```

```
kubectll exec client -- curl server
```


Architecture Flow - Issued the creation of the Service Object

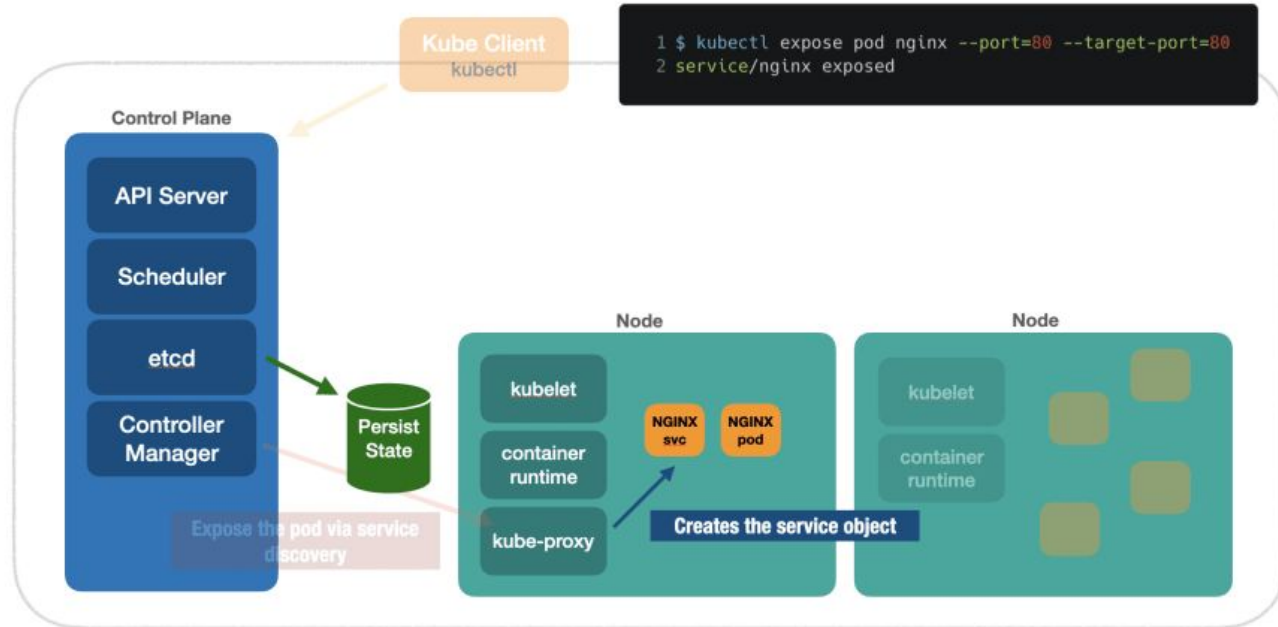


Architecture Flow - Expose the Pod via service discovery



Afterwards

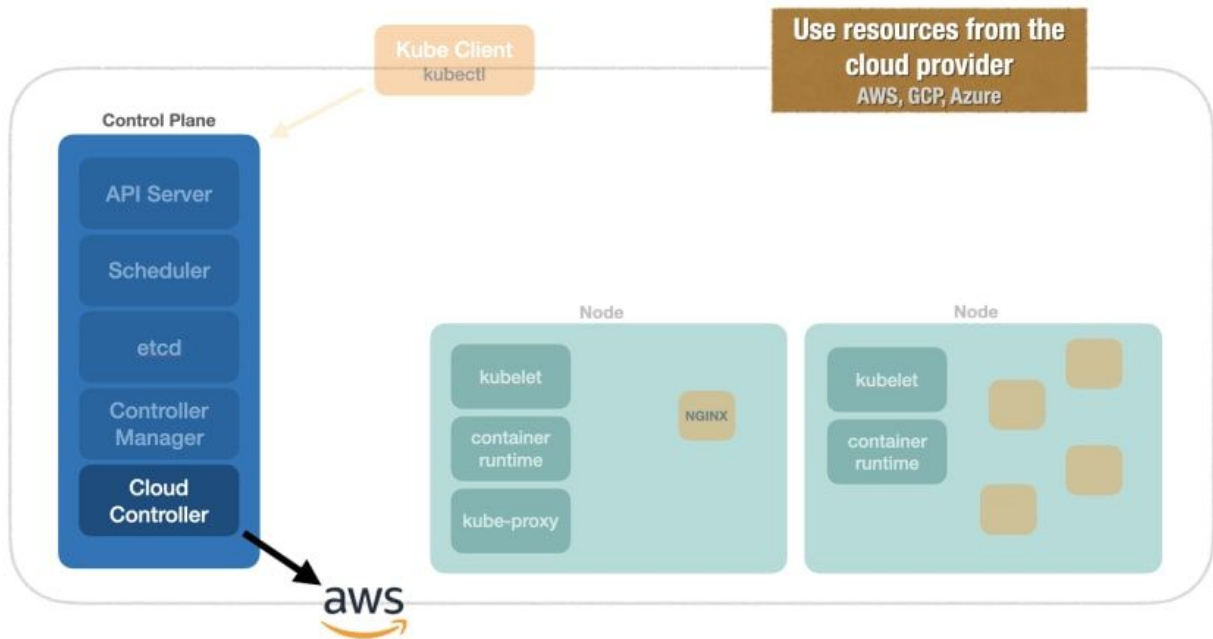
the controller manager *routes* to the **kube-proxy** component that is running in the node, which will **create the Service object** for the respective Pod. At the end of the process, the state is persisted in etcd.



Cloud Controller

Responsible for receiving requests to create objects and interacting with the underlying cloud provider if needed.

For example, when we create a Service object of type `LoadBalancer`, the Cloud Controller *will create a LB* in the underlying provider, be it AWS, GCP, Azure etc



Main Architecture Flow

