

Kubernetes 101 Tutorial - 1

05/21/2023

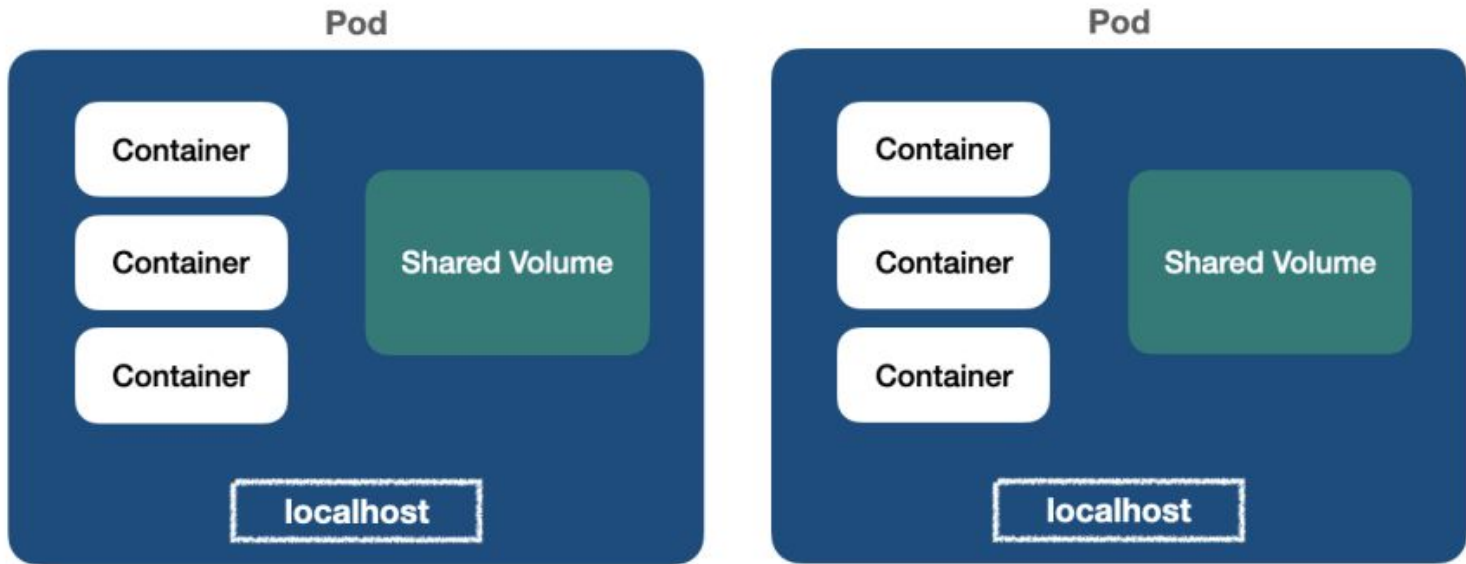
A wrapper for containers

In Kubernetes, we aren't able to create single containers directly. Instead, for the better, we can wrap containers into a single unit which comprises:

- **a specification:** where multiple containers can use the same specification as deployable units
- **a shared storage:** they can use a shared storage so the same volumes are mounted across multiple containers
- **a single network:** containers under the same wrapper can share a single network, so they can communicate to each other

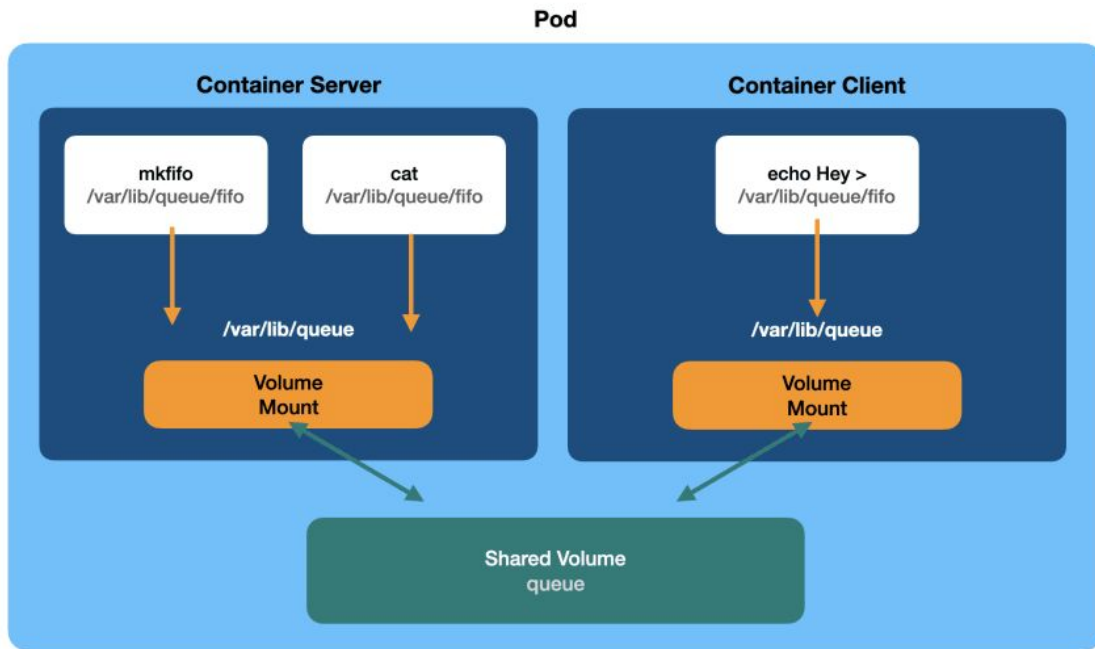

Pods

the **smallest deployable unit** you can create and manage in Kubernetes. Within Pods, we can group multiple containers that should communicate to each other somehow, either using the **same network** or through **shared volumes**.



Creating a pod

```
kubectll run <container-name> --image=<some_image>
```




YAML

Introduction: [YAML tutorial: Get started in 5 minutes \(educative.io\)](https://educative.io/yaml/tutorial)

Pod Lifecycle - Pending


It's when a Pod is **accepted by the cluster** but its *containers are not ready yet*. The Pod is **NOT yet scheduled** to any Node.



Pod Lifecycle - Running

All **containers are created** and the **Pod has been scheduled** to a Node.

At least one of the containers are still running or being started.




Pod Lifecycle - Terminated / Completed

Indicates that all the containers are `terminated` (internally by Kubernetes) or `completed`.

Pod Lifecycle - Succeeded / Failed

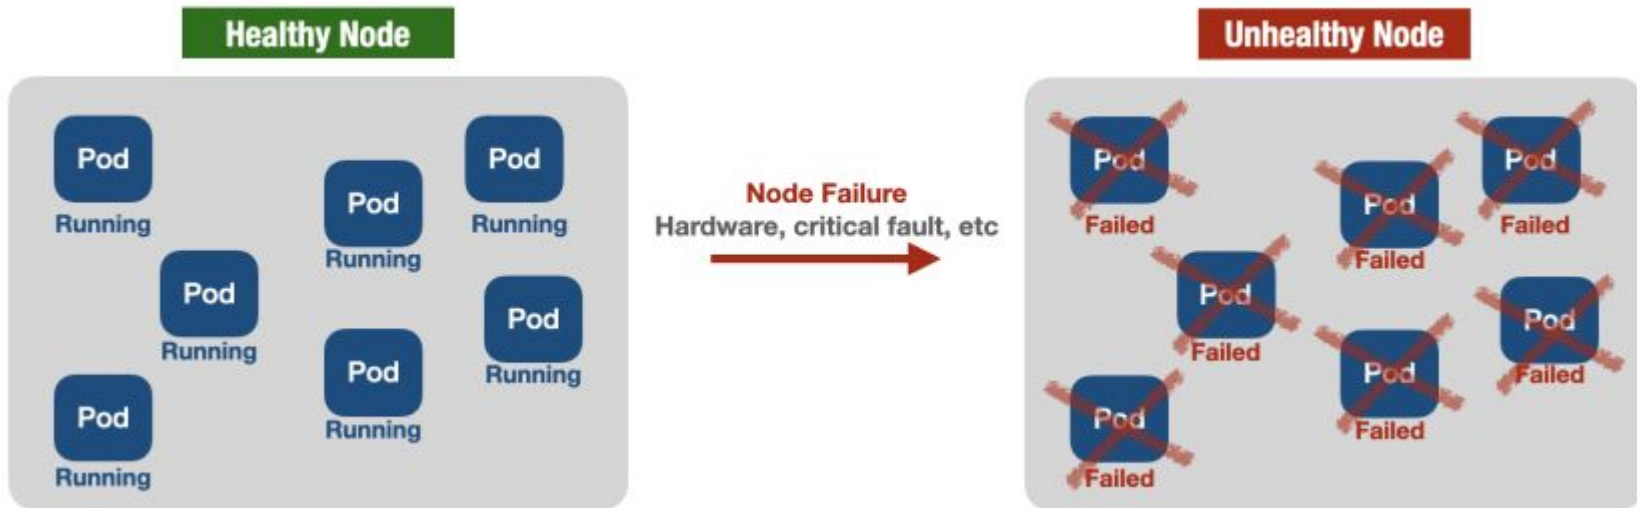
If all containers are **Terminated** in success, then the Pod status is **Succeeded**.

But in case all containers have terminated but at least 1 container terminated in failure, the Pod status is **Failed**.



Self healing

[Kubernetes Node](#): is represented by a virtual machine




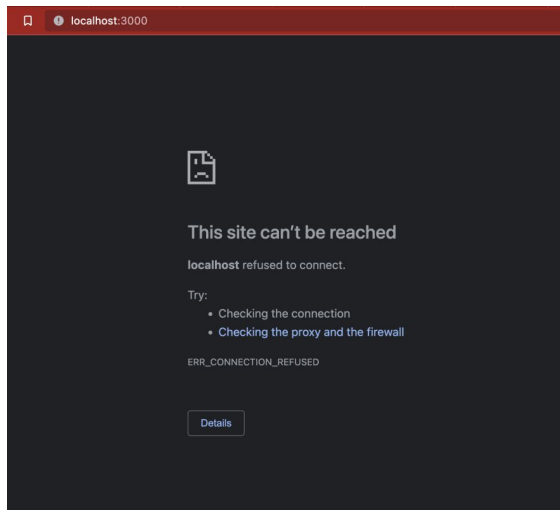
Application Downtime

A new node is required. But provisioning hardware is a *costly* operation, **it takes time**.

Meanwhile, the **Pod remains failed** in the unhealthy node and the application is suffering a **downtime**.

Once the new node *has joined and is ready to accept* new Pods, we can start **all the pods manually** using `kubectl` in the newly healthy node

```
$ kubectl apply -f ./pod.yml
```

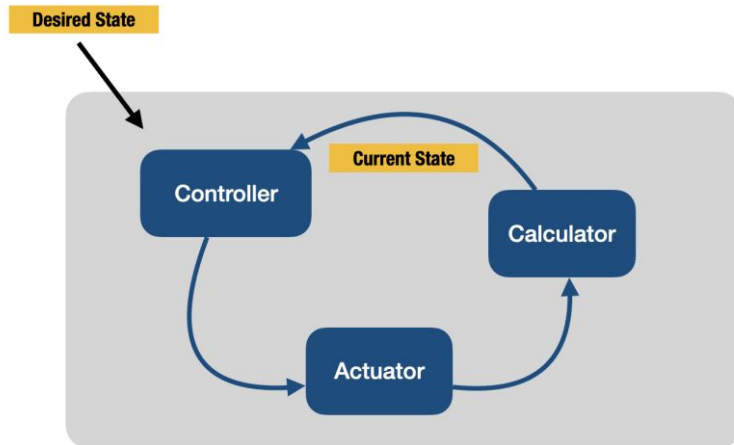


Self-healing system

A system which is capable of **detecting failures** and also *restarting components* or applications **automatically** with no human intervention.

Automation is key. And a potential solution for self-healing comes from **Robotics**.

In Robotics, we usually create a **controller** that gets a **desired state** and, by using some sort of **control loop**, it continuously check if the *current state matches the desired state*, trying to come **closer as much as possible**.



Example: a thermostat

Kubernetes Controllers

Kubernetes controllers are **control loops** that watch the cluster state, then take actions to *match the desired state* as much as possible.

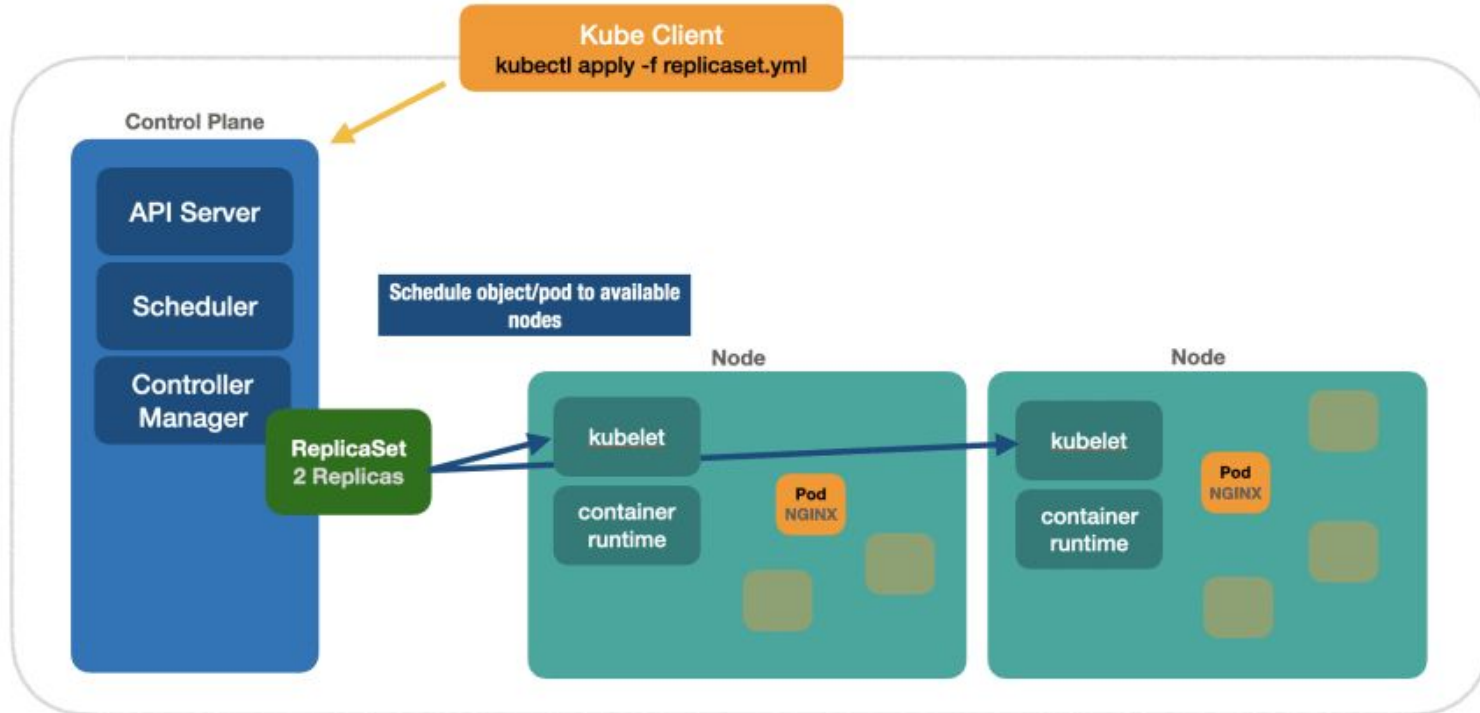
But how do we make use of *controllers*? Kubernetes provides several [Workload Resources](#) so we can rely on them to **manage Pods on our behalf**.

Time to explore one of the main *workload resources* that guarantees *self-healing capabilities*, the **ReplicaSet**.

ReplicaSet

```
1  ### The kind of the Kubernetes object
2  kind: ReplicaSet
3  apiVersion: apps/v1
4  metadata:
5  |   name: nginx
6  spec:
7  |   ### The number of replicas of nginx Pod
8  |   ### The controller will manage the Pods on our behalf
9  |   ### Anytime a Pod goes down, the controller will restart a new one to guarantee that at least 2 nginx Pods are running
10 |   replicas: 2
11 |   selector:
12 |     matchLabels:
13 |       app: nginx
14 |   template:
15 |     metadata:
16 |       labels:
17 |         app: nginx
18 |     spec:
19 |       containers:
20 |         - name: nginx
21 |           image: nginx
```

ReplicaSet



Deleting a pod of ReplicaSet/Deleting ReplicaSet

In case we delete a Pod that was created by a ReplicaSet, the controller will start a **new one automatically**.

But in case we want to delete all Pods of a ReplicaSet, we should delete the `replicaset` instead.

Next

Kubernetes workload resources